

Day 3: Tokenization in Urdu & Pashto NLP

Sami Uddin Shinwari

Subject Specialist – IT | NLP Researcher
MS Data Science (FAST NU Peshawar)
BS Computer Science (UET Peshawar)

January 15, 2026

What is Tokenization?

- Tokenization is the process of breaking text into small units called **tokens**
- Tokens can be words, subwords, characters, phrases, or numbers
- It is the **first step** of every NLP pipeline
- Tokenization makes text understandable for machines

Why Tokenization is Tricky

- Token definitions depend on language and use case
- Punctuation can split words incorrectly
- Example:

She isn't at the house.

Naive tokens: **isn** | **t**

Issue: Apostrophes break meaningful words.

Why Whitespace Works in English

- English uses spaces consistently
- Word boundaries are usually clear

This is a sentence

Tokens: This | is | a | sentence

Why Whitespace Fails in Urdu

- Space is not a reliable word boundary
- Writers may omit or add spaces

Urdu Examples:

غلط لفاظ (extra space) (no space) غلط لفاظ

Problem: Both confuse tokenizers.

Urdu Sentence Example

یہ ایک جملہ ہے

Whitespace Tokens:

ہے | جملہ | ایک | یہ

Issue: Compound words break easily in real data.

Pashto Tokenization Challenges

- Arabic-based script like Urdu
- Inconsistent spacing in informal text
- Low-resource language → limited tools

Pashto Example:

دایو هېټکه ده

Correct Form:

دا یوه بېټکه ده

Tokenization vs Chunking

- **Tokenization:** Splits text into words, subwords, or characters
- **Chunking:** Splits large text into manageable segments
- Chunking usually happens **before tokenization**
- LLMs tokenize only the most relevant chunks

Types of Tokenization

- Character Tokenization
- Word Tokenization
- Phrase Tokenization
- Sentence Tokenization
- Subword Tokenization (BPE, SentencePiece)
- Number Tokenization

Types of Tokenization (Visual)

Types of Tokenization in NLP



Why Tokenization Matters

- Wrong tokens → wrong POS tags
- NER performance drops
- TF-IDF and embeddings become noisy
- Tokenization errors propagate to all NLP tasks

Professional Uses of Tokenization

- Information retrieval (search engines)
- Text preparation and feature extraction
- Sentiment analysis
- Generative AI and chatbots
- User feedback analysis

Tokenization using NLTK (Python)

Word Tokenization Example

```
import nltk  
from nltk.tokenize import word_tokenize  
  
text = "Natural Language Processing is amazing!"  
tokens = word_tokenize(text)  
print(tokens)
```

Output:

```
['Natural', 'Language', 'Processing', 'is', 'amazing', '!']
```

Tokenization using spaCy

```
import spacy

nlp = spacy.load("en_core_web_sm")
doc = nlp("Tokenization is the first NLP step.")

for token in doc:
    print(token.text)
```

Key Feature: Context-aware tokenization

Tokenization using Hugging Face

```
from transformers import AutoTokenizer  
  
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")  
tokens = tokenizer.tokenize("Tokenization handles OOV words")  
  
print(tokens)
```

Uses: BERT, GPT, T5, Transformer models

Why Subword Tokenization?

- Handles Out-Of-Vocabulary (OOV) words
- Reduces vocabulary size
- Preserves word structure
- Used in modern NLP models

Example: *unusual* → un | usual

Why Tokenization Matters

- Affects POS tagging accuracy
- Impacts NER and embeddings
- Poor tokenization = poor NLP results

Key Takeaways

- Whitespace tokenization is NOT enough
- Urdu & Pashto need language-aware tokenizers
- Subword tokenization works best for low-resource languages
- Tokenization quality defines NLP success

#Day3 #Tokenization #UrduNLP #PashtoNLP

References

- Rehman et al. (2011). Challenges in Urdu Tokenization. ACL.
- Becker & Riaz (2012). Urdu Morphology Study.
- Khan et al. (2023). Pashto Text Processing Challenges.